

# CS5720 Neural Networks & Deep Learning ASSIGNMENT –6

NAME : MANDHA SAIKUMAR REDDY

STUDENT ID : 700765227

GITHUB LINK:

<https://github.com/saikumarreddyMandha/Assignment-4/blob/main/ICP4.ipynb>

VIDEO LINK :

[https://drive.google.com/file/d/1f4\\_mIKfWs\\_5bXu-8AN9sALISxRS6Xbl/view?usp=drive\\_link](https://drive.google.com/file/d/1f4_mIKfWs_5bXu-8AN9sALISxRS6Xbl/view?usp=drive_link)

1. Use the use case in the class:

- a. Add more Dense layers to the existing code and check how the accuracy changes.
- b. Change the data source to Breast Cancer dataset \* available in the source code folder and make required changes. Report accuracy of the model.
- c. Normalize the data before feeding the data to the model and check how the normalization change your accuracy (code given below).  
`from sklearn.preprocessing import StandardScaler  
sc = StandardScaler()`
- d. Breast Cancer dataset is designated to predict if a patient has Malignant (M) or Benign = B

## cancercancer

```
✓ 0s [2] #read the data
import pandas as pd
data = pd.read_csv('diabetes.csv')
```

```
✓ 0s [3] path_to_csv = 'diabetes.csv'
```

```
✓ 18s [4] import keras
import pandas
from keras.models import Sequential
from keras.layers import Dense, Activation

# load dataset
from sklearn.model_selection import train_test_split
import pandas as pd
import numpy as np

dataset = pd.read_csv(path_to_csv, header=None).values

X_train, X_test, Y_train, Y_test = train_test_split(dataset[:,0:8], dataset[:,8],
                                                    test_size=0.25, random_state=87)

np.random.seed(155)
my_first_nn = Sequential() # create model
my_first_nn.add(Dense(20, input_dim=8, activation='relu')) # hidden layer
my_first_nn.add(Dense(4, activation='relu')) # hidden layer
my_first_nn.add(Dense(1, activation='sigmoid')) # output layer
my_first_nn.compile(loss='binary_crossentropy', optimizer='adam', metrics=['acc'])
my_first_nn_fitted = my_first_nn.fit(X_train, Y_train, epochs=100,
                                      initial_epoch=0)

print(my_first_nn.summary())
print(my_first_nn.evaluate(X_test, Y_test))
```

```
✓ 18s ▶ Epoch 1/100
18/18 [=====] - 1s 2ms/step - loss: 5.1130 - acc: 0.5972
☐ Epoch 2/100
18/18 [=====] - 0s 2ms/step - loss: 2.1959 - acc: 0.6024
Epoch 3/100
18/18 [=====] - 0s 2ms/step - loss: 1.0648 - acc: 0.6059
Epoch 4/100
18/18 [=====] - 0s 3ms/step - loss: 0.8923 - acc: 0.5938
Epoch 5/100
18/18 [=====] - 0s 2ms/step - loss: 0.8121 - acc: 0.5955
Epoch 6/100
18/18 [=====] - 0s 2ms/step - loss: 0.7410 - acc: 0.6146
Epoch 7/100
18/18 [=====] - 0s 3ms/step - loss: 0.7091 - acc: 0.6198
Epoch 8/100
18/18 [=====] - 0s 2ms/step - loss: 0.6817 - acc: 0.6285
Epoch 9/100
18/18 [=====] - 0s 2ms/step - loss: 0.6785 - acc: 0.6128
Epoch 10/100
18/18 [=====] - 0s 2ms/step - loss: 0.6582 - acc: 0.6441
Epoch 11/100
18/18 [=====] - 0s 2ms/step - loss: 0.6489 - acc: 0.6424
Epoch 12/100
18/18 [=====] - 0s 2ms/step - loss: 0.6647 - acc: 0.6424
Epoch 13/100
18/18 [=====] - 0s 2ms/step - loss: 0.6523 - acc: 0.6632
Epoch 14/100
18/18 [=====] - 0s 2ms/step - loss: 0.6244 - acc: 0.6736
Epoch 15/100
18/18 [=====] - 0s 2ms/step - loss: 0.6247 - acc: 0.6771
Epoch 16/100
18/18 [=====] - 0s 3ms/step - loss: 0.6252 - acc: 0.6736
Epoch 17/100
18/18 [=====] - 0s 2ms/step - loss: 0.6112 - acc: 0.6771
Epoch 18/100
```

```
✓ 18s ▶ Epoch 19/100
18/18 [=====] - 0s 2ms/step - loss: 0.6103 - acc: 0.6840
↳ Epoch 20/100
18/18 [=====] - 0s 2ms/step - loss: 0.5969 - acc: 0.6892
Epoch 21/100
18/18 [=====] - 0s 2ms/step - loss: 0.5948 - acc: 0.6858
Epoch 22/100
18/18 [=====] - 0s 2ms/step - loss: 0.6157 - acc: 0.6806
Epoch 23/100
18/18 [=====] - 0s 2ms/step - loss: 0.6200 - acc: 0.6684
Epoch 24/100
18/18 [=====] - 0s 2ms/step - loss: 0.6101 - acc: 0.6840
Epoch 25/100
18/18 [=====] - 0s 2ms/step - loss: 0.5914 - acc: 0.6892
Epoch 26/100
18/18 [=====] - 0s 2ms/step - loss: 0.5902 - acc: 0.6840
Epoch 27/100
18/18 [=====] - 0s 2ms/step - loss: 0.5819 - acc: 0.6962
Epoch 28/100
18/18 [=====] - 0s 2ms/step - loss: 0.5884 - acc: 0.6910
Epoch 29/100
18/18 [=====] - 0s 2ms/step - loss: 0.5832 - acc: 0.6892
Epoch 30/100
18/18 [=====] - 0s 3ms/step - loss: 0.5827 - acc: 0.6962
Epoch 31/100
18/18 [=====] - 0s 2ms/step - loss: 0.5787 - acc: 0.6944
Epoch 32/100
18/18 [=====] - 0s 2ms/step - loss: 0.5750 - acc: 0.6997
Epoch 33/100
18/18 [=====] - 0s 2ms/step - loss: 0.5860 - acc: 0.6771
Epoch 34/100
18/18 [=====] - 0s 2ms/step - loss: 0.5857 - acc: 0.6823
Epoch 35/100
18/18 [=====] - 0s 2ms/step - loss: 0.5734 - acc: 0.6910
Epoch 36/100
```

```
✓ 18s ▶ Epoch 93/100
18/18 [=====] - 0s 3ms/step - loss: 0.5417 - acc: 0.7083
↳ Epoch 94/100
18/18 [=====] - 0s 4ms/step - loss: 0.5462 - acc: 0.7101
Epoch 95/100
18/18 [=====] - 0s 3ms/step - loss: 0.5420 - acc: 0.7153
Epoch 96/100
18/18 [=====] - 0s 2ms/step - loss: 0.5391 - acc: 0.7188
Epoch 97/100
18/18 [=====] - 0s 3ms/step - loss: 0.5370 - acc: 0.7135
Epoch 98/100
18/18 [=====] - 0s 3ms/step - loss: 0.5363 - acc: 0.7101
Epoch 99/100
18/18 [=====] - 0s 3ms/step - loss: 0.5425 - acc: 0.7153
Epoch 100/100
18/18 [=====] - 0s 3ms/step - loss: 0.5425 - acc: 0.7188
Model: "sequential"
```

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 20)	180
dense_1 (Dense)	(None, 4)	84
dense_2 (Dense)	(None, 1)	5

```
=====
Total params: 269 (1.05 KB)
Trainable params: 269 (1.05 KB)
Non-trainable params: 0 (0.00 Byte)
```

```
None
6/6 [=====] - 0s 3ms/step - loss: 0.5981 - acc: 0.6771
[0.598054826259613, 0.6770833134651184]
```

```
✓ 0s [5] #read the data
      data = pd.read_csv('breastcancer.csv')
```

```
✓ 0s [6] path_to_csv = 'sample_data/breastcancer.csv'
```

```
✓ 6s ▶ import keras (module) pd
      import pandas as pd
      import numpy as np
      from keras.models import Sequential
      from keras.layers import Dense, Activation
      from sklearn.datasets import load_breast_cancer
      from sklearn.model_selection import train_test_split

      # load dataset
      cancer_data = load_breast_cancer()
      X_train, X_test, Y_train, Y_test = train_test_split(cancer_data.data, cancer_data.target,
                                                         test_size=0.25, random_state=87)

      np.random.seed(155)
      my_nn = Sequential() # create model
      my_nn.add(Dense(20, input_dim=30, activation='relu')) # hidden layer 1
      my_nn.add(Dense(1, activation='sigmoid')) # output layer
      my_nn.compile(loss='binary_crossentropy', optimizer='adam', metrics=['acc'])
      my_nn_fitted = my_nn.fit(X_train, Y_train, epochs=100,
                               initial_epoch=0)

      print(my_nn.summary())
      print(my_nn.evaluate(X_test, Y_test))
```

```
✓ 6s ▶ 14/14 [=====] - 0s 2ms/step - loss: 0.1872 - acc: 0.9202
      Epoch 96/100
      14/14 [=====] - 0s 2ms/step - loss: 0.2868 - acc: 0.9178
      Epoch 97/100
      14/14 [=====] - 0s 2ms/step - loss: 0.1735 - acc: 0.9390
      Epoch 98/100
      14/14 [=====] - 0s 3ms/step - loss: 0.1446 - acc: 0.9249
      Epoch 99/100
      14/14 [=====] - 0s 2ms/step - loss: 0.1459 - acc: 0.9437
      Epoch 100/100
      14/14 [=====] - 0s 2ms/step - loss: 0.1356 - acc: 0.9460
      Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 20)	620
dense_4 (Dense)	(None, 1)	21

```
=====
Total params: 641 (2.50 KB)
Trainable params: 641 (2.50 KB)
Non-trainable params: 0 (0.00 Byte)
```

```
None
5/5 [=====] - 0s 5ms/step - loss: 0.2368 - acc: 0.9091
[0.2367641180753708, 0.9090909361839294]
```

```
✓ 0s [8] #read the data
      data = pd.read_csv('breastcancer.csv')
```

```
✓ 0s [9] path_to_csv = 'breastcancer.csv'
```

```
✓ 0s [10] from sklearn.preprocessing import StandardScaler
      sc = StandardScaler()
```

```
✓ 6s ▶ import keras
      import pandas as pd
      import numpy as np
      from keras.models import Sequential
      from keras.layers import Dense, Activation
      from sklearn.datasets import load_breast_cancer
      from sklearn.model_selection import train_test_split

      # load dataset
      cancer_data = load_breast_cancer()
      X_train, X_test, Y_train, Y_test = train_test_split(cancer_data.data, cancer_data.target,
                                                         test_size=0.25, random_state=87)

      np.random.seed(155)
      my_nn = Sequential() # create model
      my_nn.add(Dense(20, input_dim=30, activation='relu')) # hidden layer 1
      my_nn.add(Dense(1, activation='sigmoid')) # output layer
      my_nn.compile(loss='binary_crossentropy', optimizer='adam', metrics=['acc'])
      my_nn_fitted = my_nn.fit(X_train, Y_train, epochs=100,
                              initial_epoch=0)

      print(my_nn.summary())
      print(my_nn.evaluate(X_test, Y_test))
```

```
✓ 6s ▶ Epoch 1/100
      14/14 [=====] - 1s 2ms/step - loss: 81.4649 - acc: 0.3803
      ↗ Epoch 2/100
      14/14 [=====] - 0s 2ms/step - loss: 48.0038 - acc: 0.3803
      Epoch 3/100
      14/14 [=====] - 0s 2ms/step - loss: 15.7359 - acc: 0.4155
      Epoch 4/100
      14/14 [=====] - 0s 3ms/step - loss: 1.8353 - acc: 0.8474
      Epoch 5/100
      14/14 [=====] - 0s 2ms/step - loss: 1.3822 - acc: 0.8991
      Epoch 6/100
      14/14 [=====] - 0s 2ms/step - loss: 0.4995 - acc: 0.9131
      Epoch 7/100
      14/14 [=====] - 0s 2ms/step - loss: 0.4021 - acc: 0.9272
      Epoch 8/100
      14/14 [=====] - 0s 2ms/step - loss: 0.3421 - acc: 0.9249
      Epoch 9/100
      14/14 [=====] - 0s 2ms/step - loss: 0.3255 - acc: 0.9178
      Epoch 10/100
      14/14 [=====] - 0s 2ms/step - loss: 0.3129 - acc: 0.9178
      Epoch 11/100
      14/14 [=====] - 0s 2ms/step - loss: 0.3079 - acc: 0.9225
      Epoch 12/100
      14/14 [=====] - 0s 2ms/step - loss: 0.3394 - acc: 0.9249
      Epoch 13/100
      14/14 [=====] - 0s 2ms/step - loss: 0.3118 - acc: 0.9178
      Epoch 14/100
      14/14 [=====] - 0s 2ms/step - loss: 0.2987 - acc: 0.9178
      Epoch 15/100
      14/14 [=====] - 0s 2ms/step - loss: 0.2943 - acc: 0.9155
      Epoch 16/100
      14/14 [=====] - 0s 2ms/step - loss: 0.2984 - acc: 0.9202
      Epoch 17/100
      14/14 [=====] - 0s 2ms/step - loss: 0.2873 - acc: 0.9225
      Epoch 18/100
```



```

✓ [11] Epoch 98/100
6s 14/14 [=====] - 0s 3ms/step - loss: 0.1983 - acc: 0.9249
Epoch 99/100
14/14 [=====] - 0s 2ms/step - loss: 0.2011 - acc: 0.9366
Epoch 100/100
14/14 [=====] - 0s 2ms/step - loss: 0.1773 - acc: 0.9249
Model: "sequential_2"

Layer (type)                 Output Shape              Param #
=====
dense_5 (Dense)              (None, 20)                620
dense_6 (Dense)              (None, 1)                  21
=====
Total params: 641 (2.50 KB)
Trainable params: 641 (2.50 KB)
Non-trainable params: 0 (0.00 Byte)

None
5/5 [=====] - 0s 3ms/step - loss: 0.3693 - acc: 0.8671
[0.36931079626083374, 0.867132842540741]

```

Use Image Classification on the hand written digits data set (mnist)

1. Plot the loss and accuracy for both training data and validation data using the history object in the source code.
2. Plot one of the images in the test data, and then do inferencing to check what is the prediction of the model on that single image.
3. We had used 2 hidden layers and Relu activation. Try to change the number of hidden layer and the activation to tanh or sigmoid and see what happens.
4. Run the same code without scaling the images and check the performance?

```

import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout
import matplotlib.pyplot as plt

# load MNIST dataset
(x_train, y_train), (x_test, y_test) = mnist.load_data()

# normalize pixel values to range [0, 1]
x_train = x_train.astype('float32') / 255
x_test = x_test.astype('float32') / 255

# convert class labels to binary class matrices
num_classes = 10
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

# create a simple neural network model
model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

# train the model and record the training history
history = model.fit(x_train.reshape(-1, 784), y_train, validation_data=(x_test.reshape(-1, 784), y_test),
                    epochs=20, batch_size=128)

```

```

# plot the training and validation accuracy and loss curves
plt.figure(figsize=(10, 5))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='lower right')

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper right')

plt.show()

```

✓ 3m Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>  
11490434/11490434 [=====] - 0s 0us/step

Epoch 1/20  
469/469 [=====] - 12s 24ms/step - loss: 0.2495 - accuracy: 0.9247 - val\_loss: 0.1061 - val\_accuracy: 0.9677

Epoch 2/20  
469/469 [=====] - 11s 22ms/step - loss: 0.0997 - accuracy: 0.9690 - val\_loss: 0.0840 - val\_accuracy: 0.9739

Epoch 3/20  
469/469 [=====] - 11s 24ms/step - loss: 0.0717 - accuracy: 0.9771 - val\_loss: 0.0751 - val\_accuracy: 0.9776

Epoch 4/20  
469/469 [=====] - 10s 22ms/step - loss: 0.0568 - accuracy: 0.9817 - val\_loss: 0.0759 - val\_accuracy: 0.9761

Epoch 5/20  
469/469 [=====] - 10s 21ms/step - loss: 0.0467 - accuracy: 0.9848 - val\_loss: 0.0679 - val\_accuracy: 0.9796

Epoch 6/20  
469/469 [=====] - 10s 21ms/step - loss: 0.0369 - accuracy: 0.9880 - val\_loss: 0.0636 - val\_accuracy: 0.9815

Epoch 7/20  
469/469 [=====] - 9s 20ms/step - loss: 0.0371 - accuracy: 0.9876 - val\_loss: 0.0707 - val\_accuracy: 0.9787

Epoch 8/20  
469/469 [=====] - 11s 23ms/step - loss: 0.0305 - accuracy: 0.9898 - val\_loss: 0.0680 - val\_accuracy: 0.9802

Epoch 9/20  
469/469 [=====] - 10s 22ms/step - loss: 0.0268 - accuracy: 0.9912 - val\_loss: 0.0752 - val\_accuracy: 0.9807

Epoch 10/20  
469/469 [=====] - 10s 22ms/step - loss: 0.0267 - accuracy: 0.9911 - val\_loss: 0.0737 - val\_accuracy: 0.9820

Epoch 11/20  
469/469 [=====] - 9s 19ms/step - loss: 0.0242 - accuracy: 0.9922 - val\_loss: 0.0664 - val\_accuracy: 0.9840

Epoch 12/20  
469/469 [=====] - 10s 21ms/step - loss: 0.0242 - accuracy: 0.9918 - val\_loss: 0.0645 - val\_accuracy: 0.9833

Epoch 13/20  
469/469 [=====] - 10s 21ms/step - loss: 0.0213 - accuracy: 0.9929 - val\_loss: 0.0693 - val\_accuracy: 0.9832

Epoch 14/20  
469/469 [=====] - 9s 20ms/step - loss: 0.0187 - accuracy: 0.9937 - val\_loss: 0.0700 - val\_accuracy: 0.9841

Epoch 15/20  
469/469 [=====] - 10s 20ms/step - loss: 0.0183 - accuracy: 0.9940 - val\_loss: 0.0966 - val\_accuracy: 0.9791


Epoch 16/20  
469/469 [=====] - 10s 22ms/step - loss: 0.0169 - accuracy: 0.9942 - val\_loss: 0.0772 - val\_accuracy: 0.9821

Epoch 17/20  
469/469 [=====] - 10s 22ms/step - loss: 0.0162 - accuracy: 0.9945 - val\_loss: 0.0756 - val\_accuracy: 0.9847

✓ 3m Epoch 18/20  
469/469 [=====] - 10s 21ms/step - loss: 0.0168 - accuracy: 0.9947 - val\_loss: 0.0728 - val\_accuracy: 0.9847

Epoch 19/20  
469/469 [=====] - 9s 20ms/step - loss: 0.0142 - accuracy: 0.9952 - val\_loss: 0.0845 - val\_accuracy: 0.9816

Epoch 20/20  
469/469 [=====] - 10s 21ms/step - loss: 0.0169 - accuracy: 0.9946 - val\_loss: 0.0837 - val\_accuracy: 0.9843





```

import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout
import matplotlib.pyplot as plt
import numpy as np

# load MNIST dataset
(x_train, y_train), (x_test, y_test) = mnist.load_data()

# normalize pixel values to range [0, 1]
x_train = x_train.astype('float32') / 255
x_test = x_test.astype('float32') / 255

# convert class labels to binary class matrices
num_classes = 10
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

# create a simple neural network model
model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])


# train the model
model.fit(x_train.reshape(-1, 784), y_train, validation_data=(x_test.reshape(-1, 784), y_test),

```

```

469/469 [=====] - 9s 20ms/step - loss: 0.0205 - accuracy: 0.9933 - val_loss: 0.0763 - val_accuracy: 0.9803
Epoch 13/20
469/469 [=====] - 10s 21ms/step - loss: 0.0217 - accuracy: 0.9926 - val_loss: 0.0759 - val_accuracy: 0.9826
Epoch 14/20
469/469 [=====] - 10s 21ms/step - loss: 0.0177 - accuracy: 0.9942 - val_loss: 0.0756 - val_accuracy: 0.9814
Epoch 15/20
469/469 [=====] - 9s 20ms/step - loss: 0.0200 - accuracy: 0.9936 - val_loss: 0.0739 - val_accuracy: 0.9822
Epoch 16/20
469/469 [=====] - 10s 20ms/step - loss: 0.0193 - accuracy: 0.9937 - val_loss: 0.0647 - val_accuracy: 0.9845
Epoch 17/20
469/469 [=====] - 10s 21ms/step - loss: 0.0162 - accuracy: 0.9947 - val_loss: 0.0811 - val_accuracy: 0.9807
Epoch 18/20
469/469 [=====] - 11s 23ms/step - loss: 0.0170 - accuracy: 0.9946 - val_loss: 0.0790 - val_accuracy: 0.9820
Epoch 19/20
469/469 [=====] - 10s 22ms/step - loss: 0.0140 - accuracy: 0.9954 - val_loss: 0.0791 - val_accuracy: 0.9808
Epoch 20/20
469/469 [=====] - 9s 19ms/step - loss: 0.0148 - accuracy: 0.9949 - val_loss: 0.0826 - val_accuracy: 0.9829

```





```
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout
import matplotlib.pyplot as plt
import numpy as np

# load MNIST dataset
(x_train, y_train), (x_test, y_test) = mnist.load_data()


# normalize pixel values to range [0, 1]
x_train = x_train.astype('float32') / 255
x_test = x_test.astype('float32') / 255

# convert class labels to binary class matrices
num_classes = 10
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)


# create a list of models to train
models = []


# model with 1 hidden layer and tanh activation
model = Sequential()
model.add(Dense(512, activation='tanh', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))
models.append(('1 hidden layer with tanh', model))

# model with 1 hidden layer and sigmoid activation
model = Sequential()
model.add(Dense(512, activation='sigmoid', input_shape=(784,)))
```




1 hidden layer with tanh - Test loss: 0.0668, Test accuracy: 0.9815





1 hidden layer with sigmoid - Test loss: 0.0622, Test accuracy: 0.9821

### 2 hidden layers with tanh



2 hidden layers with tanh - Test loss: 0.0704, Test accuracy: 0.9813

```

import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout
import matplotlib.pyplot as plt
import numpy as np

# load MNIST dataset
(x_train, y_train), (x_test, y_test) = mnist.load_data()

# convert class labels to binary class matrices
num_classes = 10
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

# create a list of models to train
models = []

# model with 1 hidden layer and tanh activation
model = Sequential()
model.add(Dense(512, activation='tanh', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))
models.append(('1 hidden layer with tanh', model))

# model with 1 hidden layer and sigmoid activation
model = Sequential()
model.add(Dense(512, activation='sigmoid', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))
models.append(('1 hidden layer with sigmoid', model))

```



```


model.add(Dense(num_classes, activation='softmax'))
models.append(('2 hidden layers with tanh', model))

# model with 2 hidden layers and sigmoid activation
model = Sequential()
model.add(Dense(512, activation='sigmoid', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='sigmoid'))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))
models.append(('2 hidden layers with sigmoid', model))

# train each model and plot loss and accuracy curves
for name, model in models:
    model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
    history = model.fit(x_train.reshape(-1, 784), y_train, validation_data=(x_test.reshape(-1, 784), y_test),
                        epochs=20, batch_size=128, verbose=0)
    # plot loss and accuracy curves
    plt.plot(history.history['loss'], label='train_loss')
    plt.plot(history.history['val_loss'], label='val_loss')
    plt.plot(history.history['accuracy'], label='train_accuracy')
    plt.plot(history.history['val_accuracy'], label='val_accuracy')
    plt.title(name)
    plt.xlabel('Epoch')
    plt.legend()
    plt.show()


# evaluate the model on test data
loss, accuracy = model.evaluate(x_test.reshape(-1, 784), y_test, verbose=0)
print('{:} - Test loss: {:.4f}, Test accuracy: {:.4f}'.format(name, loss, accuracy))

```




1 hidden layer with tanh - Test loss: 0.1909, Test accuracy: 0.9436


1 hidden layer with sigmoid



1 hidden layer with sigmoid - Test loss: 0.1449, Test accuracy: 0.9580

2 hidden layers with tanh





2 hidden layers with sigmoid - Test loss: 0.1039, Test accuracy: 0.9652